

**POSTGRE(NO)SQL**

# WHY USING JSON

Flexible schema (≠ destructured)

Open Library, Wikidata, ...

Data denormalization

Avoid complex joins and reduce number of queries  
(e.g. EAV model)

Native JSON documents storage

JSON responses from REST APIs

# THE OLD WAY

Before NoSQL and native JSON

# EAV (ENTITY-ATTRIBUTE-VALUE) MODEL

One table contains the entities, another table contains the names of the properties (attributes) and a third table links the entities with their attributes and holds the value. This gives you the flexibility for having **different sets of properties** [...].

Selecting one or more entities based on 1 attribute value requires 2 **joins** [...]. Also, the properties usually are all stored as strings, which results in **type casting**, both for the result as for the WHERE clause.




JSONB has potential for greatly **simplifying schema design without sacrificing query performance**.

Replacing EAV with JSONB in PostgreSQL

via Wikipedia / Entity-attribute-value model

# PLAIN TEXT AND CLIENT-SIDE MAPPING

Store JSON data in a text field and process it client-side.

-  No SQL querying features
-  Mapping overhead on client-side
-  Multiple client bring duplicated code

# **THE MODERN WAY**

2013 and beyond

# 2012-09-10

## PostgreSQL 9.2 introduces j s o n data type

- Store an exact copy of the input
  - Preserve order of keys
  - Preserve duplicate key/value pairs
- JSON validation (according to [RFC-7159](#))
- JSON-specific functions and operators
  - Must *reparse input* on each execution
  - If duplicate key/value pairs are present, only the last value is considered

# 2 YEARS LATER (2014-12-18)

## PostgreSQL 9.4 introduces jsonb data type

- Slower to input due to parsing overhead
  - Does not preserve the order of object keys
  - Does not keep duplicate object keys (in case of duplicate keys, only last value is kept)
- Faster to process (no need to reparse input)
- Support indexing
- Use PostgreSQL types to map JSON primitive types



## 1 YEAR LATER (2015)

3rd August: *MySQL 5.7.8* introduces **JSON** data type

14th October: *SQLite 3.9.0* introduces **JSON** data type

# JSONB OPERATIONS

SELECT



UPDATE

QUERY

```
CREATE TABLE authors (  
    type character varying,  
    key character varying NOT NULL PRIMARY KEY,  
    revision integer,  
    last_modified timestamp without time zone,  
    json jsonb  
);  
  
CREATE INDEX authors_idx_gin ON authors USING gin (json);  
CREATE INDEX authors_idx_ginp ON authors USING gin (  
    json jsonb_path_ops  
);
```

```
{
  "key": "/authors/OL2623297A",
  "name": "Sir Arthur Conan Doyle",
  "photos": [5541405],
  "created": {
    "type": "/type/datetime", "value": "2008-04-29T13:35:46.876380"
  },
  "remote_ids": {
    "wikidata": "Q35610"
  },
  "last_modified": {
    "type": "/type/datetime", "value": "2017-03-31T14:21:52.424685"
  }
}
```

# SELECT

SELECT BY  \ RESULT TYPE 	jsonb	text
<i>Object key/array index</i>	->	->>
<i>JSON path</i>	#>	#>>

## OBJECT KEY

```
SELECT json->'name', json->>'name', -- By object key
       json #> '{name}', json #>> '{name}' -- By JSON path
FROM authors;
```

```
SELECT json->'remote_ids'->'wikidata',
       json->'remote_ids'->>'wikidata',
       json #> '{remote_ids,wikidata}',
       json #>> '{remote_ids,wikidata}'
FROM authors;
```

If a field doesn't exist, **NULL** is returned.

## ARRAY ELEMENT

```
-- Positive index (starting from 0)
SELECT json->'photos'->0, json->'photos'->>0, -- By index
       json #> '{photos, 0}', json #>> '{photos, 0}' -- By JSON path
FROM authors;
```

```
-- Negative indices count from the end of the array
SELECT json->'photos'->-1, json->'photos'->>-1, -- By index
       json #> '{photos, -1}', json #>> '{photos, -1}' -- By JSON path
FROM authors;
```

If an array element doesn't exist, **NULL** is returned.

# UPDATE

Concatenate

Delete

`jsonb_strip_nulls`

`jsonb_set`



## CONCATENATE ( | | )

```
-- ARRAY
SELECT '["EUR", "USD"]'::jsonb || '"GBP"'::jsonb;
-- '["EUR", "USD", "GBP"]'
SELECT '["EUR", "USD"]'::jsonb || '{"GBP": "Pound Sterling"}'::jsonb;
-- '["EUR", "USD", {"GBP": "Pound Sterling"}]'
SELECT '["EUR", "USD"]'::jsonb || '["GBP", "CAD"]'::jsonb;
-- '["EUR", "USD", "GBP", "CAD"]'
```

```
-- OBJECTS
-- If duplicate keys are found, only last value is kept
SELECT '{"EUR": "Euro", "USD": "United States Dollar"}'::jsonb ||
      '{"GBP": "Pound Sterling", "USD": "US Dollar"}'::jsonb;
-- '{"EUR": "Euro", "GBP": "Pound Sterling", "USD": "US Dollar"}'
```

## DELETE (-, #-)

```
-- OBJECT FIELD
SELECT '{"EUR": "Euro", "USD": "United States Dollar",
      "GBP": "Pound Sterling"}'::jsonb - 'USD', -- By key name
      '{"EUR": "Euro", "USD": "United States Dollar",
      "GBP": "Pound Sterling"}'::jsonb #- '{"USD"}'; -- By key path
-- '{"EUR": "Euro", "GBP": "Pound Sterling"}'
```

```
-- ARRAY ELEMENT
SELECT '["EUR", "USD", "GBP"]'::jsonb - 1, -- '["EUR", "GBP"]'
      '["EUR", "USD", "GBP"]'::jsonb #- '{1}', -- '["EUR", "GBP"]'
      '["EUR", "USD", "GBP"]'::jsonb - -1, -- '["EUR", "USD"]'
      '["EUR", "USD", "GBP"]'::jsonb #- '{-1}'; -- '["EUR", "USD"]'
```

## JSONB\_STRIP\_NULLS(FROM\_JSON)

```
SELECT jsonb_strip_nulls('{"EUR": "Euro", "GBP": null}'::jsonb);  
-- '{"EUR": "Euro"}'
```

Remove object fields with null value.

```
SELECT jsonb_strip_nulls('["EUR", null, null, "CAD"]'::jsonb);  
-- '["EUR", null, null, "CAD"]'
```

Has no effects on arrays with null elements.

## JSONB\_SET(TARGET, PATH, NEW\_VALUE[, CREATE\_IF\_MISSING])

```
SELECT jsonb_set(  
    '{"code": "EUR", "name": "Euro", "symbol": null}', '{symbol}', '€'  
); -- '{"code": "EUR", "name": "Euro", "symbol": "€"}'
```

```
SELECT jsonb_set(  
    '{"code": "EUR", "name": "Euro"}', '{symbol}', '€'  
); -- '{"code": "EUR", "name": "Euro", "symbol": "€"}'  
SELECT jsonb_set(  
    '{"code": "EUR", "name": "Euro"}', '{symbol}', '€', false  
); -- '{"code": "EUR", "name": "Euro"}'
```

Update/insert a JSON value at given object path and/or array position.

# QUERY

Existence

Containment

Expansion

# EXISTENCE (?, ? |, ?&)

## -- ARRAY ELEMENTS

```
SELECT '["EUR", "USD", "GBP"]'::jsonb ? 'EUR'; -- true
SELECT '["EUR", "USD", "GBP"]'::jsonb ?| ARRAY['EUR', 'CAD']; -- true
SELECT '["EUR", "USD", "GBP"]'::jsonb ?& ARRAY['EUR', 'CAD']; -- false
-- Work only with elements on top-level array
SELECT '["EUR", ["USD", "GBP"]]'::jsonb ? 'USD'; -- false
```

## -- OBJECT KEYS

```
SELECT '{"EUR": "Euro", "USD": "United States Dollar", "GBP":  
      "Pound Sterling"}'::jsonb ? 'EUR'; -- true
SELECT '{"EUR": "Euro", "USD": "United States Dollar", "GBP":  
      "Pound Sterling"}'::jsonb ?| ARRAY['EUR', 'CAD']; -- true
SELECT '{"EUR": "Euro", "USD": "United States Dollar", "GBP":  
      "Pound Sterling"}'::jsonb ?& ARRAY['EUR', 'CAD']; -- false
-- Work only with keys on top-level object
SELECT '{"currencies": {"EUR": "Euro", "USD": "United States Dollar",  
      "GBP": "Pound Sterling"}}'::jsonb ? 'EUR'; -- false
```

Check if a string appears as an object key or array element at the top level.

# CONTAINMENT (@>, <@)

```
-- PRIMITIVE TYPES
-- Check absolute equality
SELECT '"EUR"'::jsonb @> '"EUR"'::jsonb; -- true
SELECT '"EUR"'::jsonb @> '"GBP"'::jsonb; -- false
```

```
-- ARRAYS
-- Order of elements is ignored
-- Duplicate elements are considered only once
SELECT '["EUR", "GBP", "USD"]'::jsonb
      @> '["GBP", "EUR", "EUR"]'::jsonb; -- true
SELECT '["EUR", "GBP", "USD"]'::jsonb @> '"EUR"'::jsonb; -- true
SELECT '["EUR", "GBP", "USD"]'::jsonb @> '"CAD"'::jsonb; -- false
```

```
-- OBJECTS
-- Structure and data contents must both match
SELECT '{"id": 12345, "name": "Nicola", "surname": "Moretto"}'::jsonb
      @> '{"name": "Nicola"}'; -- true
SELECT '{"id": 12345, "name": "Nicola", "surname": "Moretto"}'::jsonb
      @> '{"name": "Francesco"}'; -- false
```

## EXPANSION (PART I: OBJECTS)

```
-- Return a record for each object field
SELECT jsonb_object_keys(
    '{"code": "EUR", "name": "Euro", "symbol": "€"}' ::jsonb
);
```

```
-- Return a record with 4 columns and values from matching object keys
SELECT *
FROM jsonb_to_record(
    '{"code":"EUR", "number":478, "name":"Euro", "symbol":"€"}' ::jsonb
) AS r(code text, number int, name text, symbol text);
```



## EXPANSION (PART II: ARRAYS)

```
-- Return a record for each array element (as jsonb or text)
SELECT jsonb_array_elements('["EUR", "GBP", "CAD", "USD"]'::jsonb),
       jsonb_array_elements_text('["EUR", "GBP", "CAD", "USD"]'::jsonb);
```

```
-- Return for each object a record with two columns (code and name)
-- and values from matching object fields
SELECT *
FROM jsonb_to_recordset('[
  {"code": "GBP", "name": "Pound Sterling", "symbol": "£"},
  {"code": "EUR", "name": "Euro", "symbol": "€"},
  {"code": "USD", "name": "United States Dollar", "symbol": "$"}
]'::jsonb) AS c(code text, name text)
ORDER BY code ASC; -- Sort records by currency ISO code
```

# JSONB INDEXING

```
CREATE INDEX ON authors USING GIN (?)
```

▼ OPERATOR	▼ TARGET	
	Table column	Index expression
<i>jsonb_ops</i>	json	(json->'name')
<i>jsonb_path_ops</i>	json jsonb_path_ops	(json->'name') jsonb_path_ops

# INDEX EXPRESSIONS

Index expression is required if JSONB operator is not applied to table column

*An index column need not be just a column of the underlying table, but can be a function or scalar expression computed from one or more columns of the table.*

*PostgreSQL 9.6 Documentation / Indexes / Indexes on Expressions*

# JSONB\_OPS (DEFAULT)

Creates an index item for each **key** and value in the data.

- 👍 Support key-exists (`?`, `?&` and `? |`) and path/value-exists operators (`@>`, `<@`)
- 👎 Slower search than `jsonb_path_ops`
- 👎 Bigger index than `jsonb_path_ops` on same data

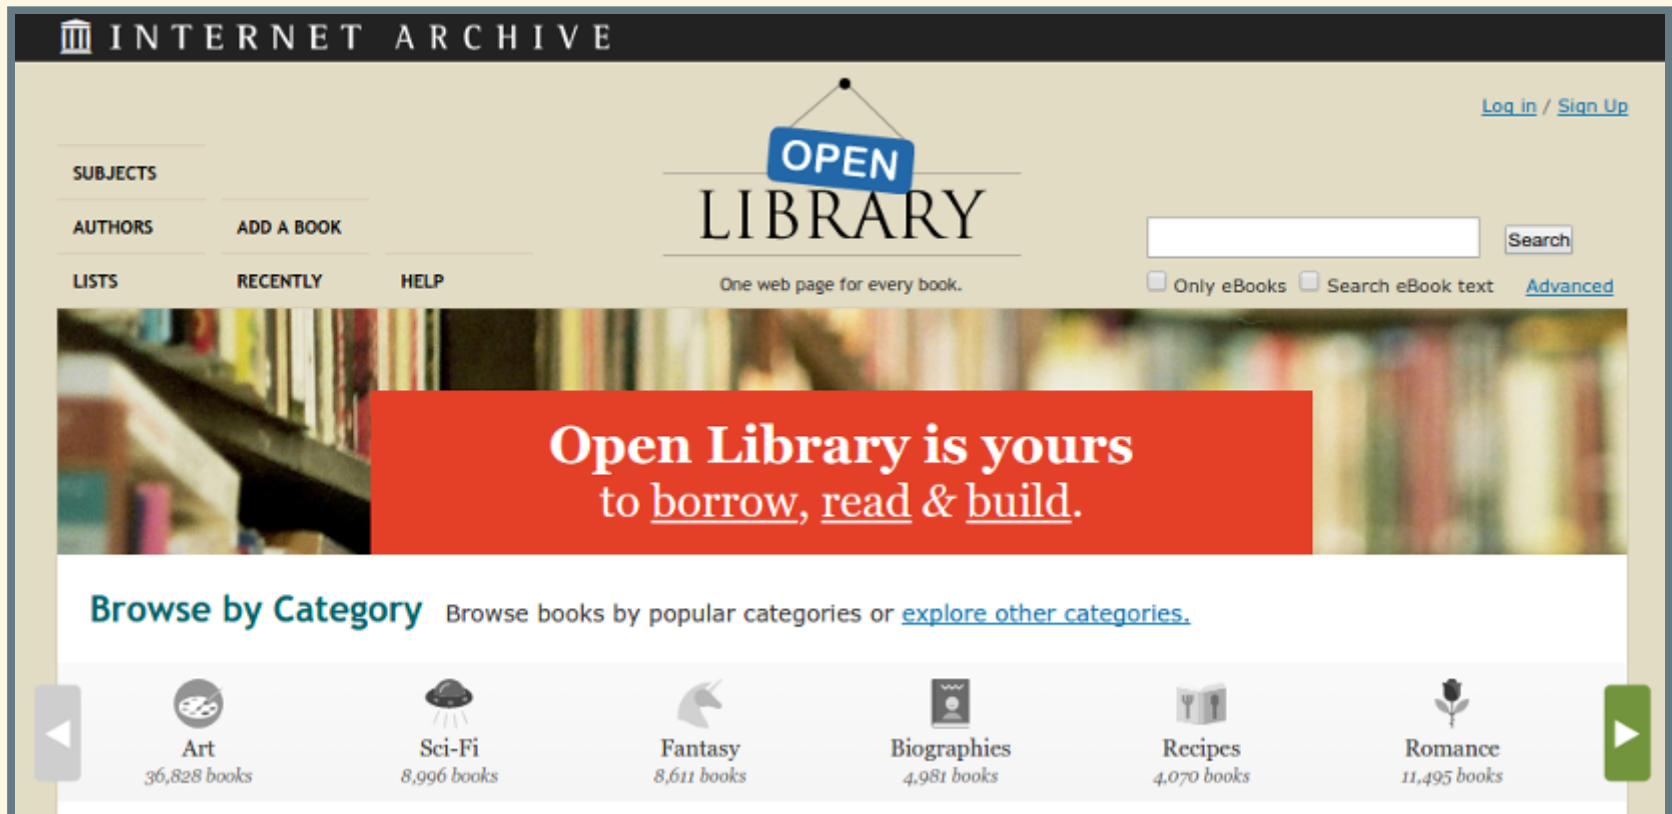
# JSONB\_PATH\_OPS

Creates an index entry for each **value** and the key(s) leading to it.

- 👍 Best for **containment queries**
- 👍 Faster search than `jsonb_ops`
- 👍 Smaller index than `jsonb_ops` on same data
- 👎 Support only path/value-exists operators (`@>`, `<@`)
- 👎 No entries for JSON structures without value

# DEMO

6.964.927 authors



# SELECTING RECORD

```
SELECT * FROM authors
WHERE key = '/authors/OL2623297A';
-- Index Scan using PRIMARY KEY on authors
```

```
SELECT * FROM authors
WHERE json->>'key' = '/authors/OL2623297A';
-- Seq Scan on authors
```

```
SELECT * FROM authors
WHERE json @> '{"key": "/authors/OL2623297A"}' ::jsonb;
-- Index Scan on authors_idx_ginp
```

```
SELECT * FROM authors
WHERE json @> '{"name": "Sir Arthur Conan Doyle"}' ::jsonb;
-- Index Scan on authors_idx_ginp
```

# COUNTING RECORDS

```
SELECT COUNT(*) FROM authors
WHERE json->'remote_ids' ? 'wikidata';
-- Seq Scan on authors
```

Existence operator can only use simple index for top-level keys, **remember?**

```
-- Create an index targeting the specific key
CREATE INDEX authors_idx_gin_remoteids
ON authors USING GIN ((json -> 'remote_ids'));
```

```
SELECT COUNT(*) FROM authors
WHERE json->'remote_ids' ? 'wikidata';
-- Index Scan on authors_idx_gin_remoteids
```



# AGGREGATING RECORDS

```
CREATE INDEX authors_idx_ginp_created
ON authors USING GIN ((json -> 'created') jsonb_path_ops);
```

```
-- Number of records added each day of last 2 months
SELECT (json #>> '{created,value}')::timestamp::date AS day,
       COUNT(*) AS records
FROM authors
WHERE json->'created' @> '{"type": "/type/datetime"}'
      AND (json #>> '{created,value}')::timestamp
          >= (CURRENT_DATE - INTERVAL '2 month')
GROUP BY 1
ORDER BY 1 DESC
-- Bitmap Index Scan on authors_idx_ginp_created
```

You can aggregate on an arbitrary date part using `date_trunc` function

```
CREATE MATERIALIZED VIEW mv_authors_cts AS
SELECT key, (json #>> '{created,value}'):timestamp AS ts
FROM authors
WHERE json->'created' @> '{"type": "/type/datetime"}'
WITH DATA;
CREATE INDEX mv_authors_cts_ts_idx ON mv_authors_cts(ts);

-- Remember to refresh materialized view if source table changes!
REFRESH MATERIALIZED VIEW mv_authors_cts;
```

```
SELECT ts::date AS day, COUNT(*) AS records
FROM mv_authors_cts
WHERE ts >= (CURRENT_DATE - INTERVAL '2 month')
GROUP BY 1
ORDER BY 1 DESC
-- Bitmap Index Scan on mv_authors_lmts_ts_idx
```


# Thank you!

Full presentation and curated links on [GitHub](#)

 Nicola Moretto

 [n.moretto@nicolamoretto.eu](mailto:n.moretto@nicolamoretto.eu)

 [nicola88 \(GitHub\)](#)

 [nicola88 \(BitBucket\)](#)